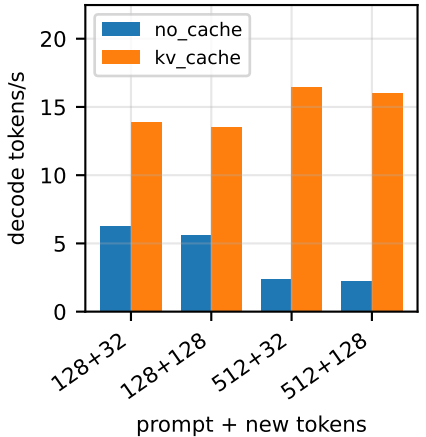
no_cache
kv_cache
decode tokens/s
0
5
10
15
20
128+32
128+128
512+32
512+128
prompt + new tokens

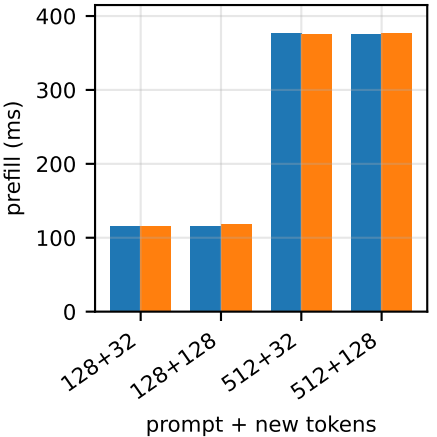
prefill (ms)
0
100
200
300
400
128+32
128+128
512+32
512+128
prompt + new tokens

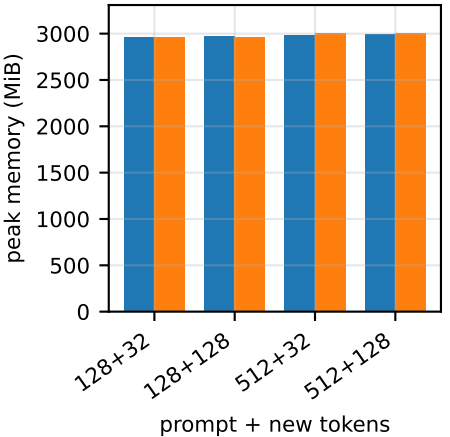
peak memory (MiB)
0
500
1000
1500
2000
2500
3000
128+32
128+128
512+32
512+128
prompt + new tokens