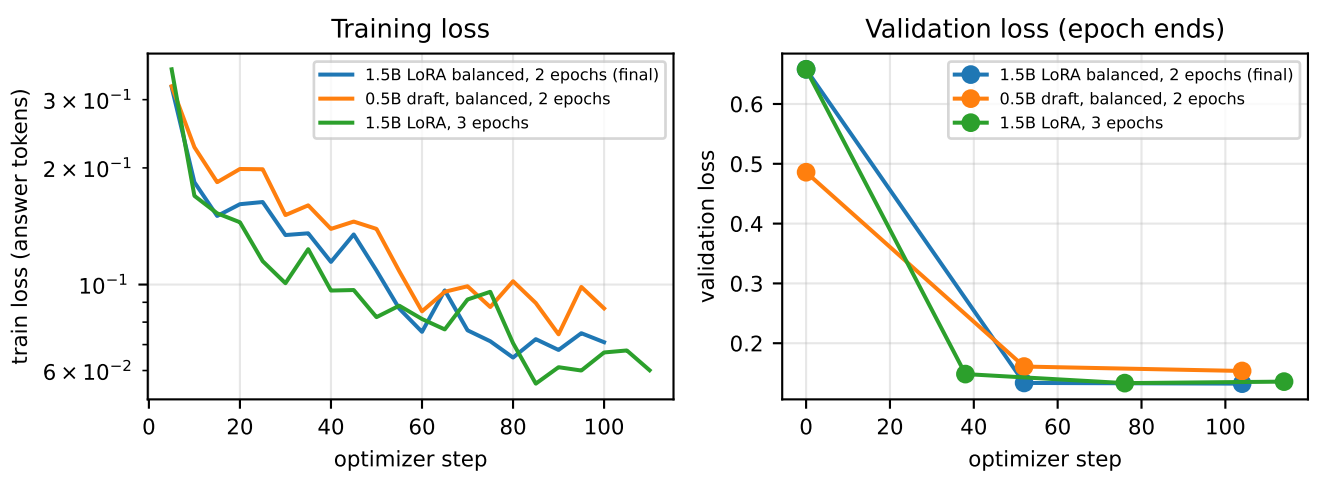
Training loss
train loss (answer tokens)
3×10^{-1}
2×10^{-1}
10^{-1}
6×10^{-2}
0
20
40
60
80
100
optimizer step
1.5B LoRA balanced, 2 epochs (final)
0.5B draft, balanced, 2 epochs
1.5B LoRA, 3 epochs
Validation loss (epoch ends)
validation loss
0.6
0.5
0.4
0.3
0.2
0
20
40
60
80
100
optimizer step
1.5B LoRA balanced, 2 epochs (final)
0.5B draft, balanced, 2 epochs
1.5B LoRA, 3 epochs